CO 485/685 Fall 2022:
Lecture Notes

Lecture notes taken, unless otherwise specified, by myself during the Fall 2022 offering of CO 485/685, taught by David Jao.

Chapter/lecture titles are made-up nonsense and do not follow the textbook or any other published resource. Actually, scratch that, this entire document is nonsense because I am literally auditing this course two nested prerequisites behind.

Chapter 1

Introduction to Cryptography

Lecture 1   (09/07; skipped)

Lecture 2   Almost-Public Key Cryptosystems (09/09)

- For a symmetric key cryptosystem, require sets of key space K, message space M, and ciphertext space C
    - Define encryption function $Enc : K \to M \to C$ and decryption $Dec : K \to C \to M$
    - Correctness property: for all k, $Dec(k)$ is a left inverse of $Enc(k)$
    - Symmetric means that both decryption and encryption use shared secret k, which we assume is drawn randomly from K
- Public key encryption scheme (Diffie, Hellman, Merkle, c. 1976)
    - Setup similar: message space M and ciphertext space C but with two key spaces K_1 of public keys and K_2 of private keys
    - Define $Enc : K_1 \to M \to C$ and $Dec : K_2 \to C \to M$
    - Define $KeyGen : \mathbb{1}^\ell \to R \subset K_1 \times K_2$
        * For some reason, let $\mathbb{1}^n$ be the unary representation of n??
    - Correctness: for all $(k_1, k_2) \in R$ related, $Dec(k_2)$ is a left inverse of $Enc(k_1)$
- Merkle puzzle (1974)
    - Each party creates "puzzle" which is hard to solve but not too hard
    - Alice generates 1,000,000 puzzles and sends them to Bob
    - Bob solves one of the puzzles arbitrarily and sends half of the answer to Alice
    - Alice knows the answer, so Alice knows the second half of the answer, which becomes the shared secret
    - Eve cannot (realistically) solve 500,000 puzzles in time to intercept
- Diffie–Hellman key exchange
    - Consider the multiplicative group $G = (\mathbb{Z}/p\mathbb{Z})^* = 1, \dots, p-1$ and some arbitrary element $g \in G$ with sufficiently large order
    - Alice privately picks some $x \in \mathbb{Z}$, computes g^x, and sends it to Bob
    - Bob privately picks some $y \in \mathbb{Z}$, computes g^y, and sends it to Alice
    - Both can now calculate a shared secret $k = g^{xy} = (g^x)^y = (g^y)^x$
    - Eve would have to solve the Diffie–Hellman problem: given p, g, g^x, g^y, find g^{xy} which is known to be hard
- Clifford Cocks privately discovered RSA 1973, DH 1974 for GCHQ (if you believe the intelligence community)

Lecture 3   A Public Key Cryptosystem – RSA (09/12)

- RSA (Rivest, Shamir, Adleman 1977): first cryptosystem and remains secure
- Theoretically secure, but implementations are ass (cf. "Fuck RSA")
- MATH 135 review of the algorithm:
    - This "textbook RSA" has practical flaws and is insecure
    - $KeyGen : \mathbb{1}^\ell \to (pk, sk) \in R$
        1. Choose random primes $p, q \approx 2^\ell$ where p and q are odd and distinct
        2. Compute $n = pq$
        3. Choose $e \in (\mathbb{Z}/\phi(n)\mathbb{Z})^\times$ where $\phi(n) = (p-1)(q-1)$
        4. Compute $d = e^{-1} \bmod \phi(n)$
        5. Disclose public key (n, e) and keep secret key (n, d)
    - $Enc : K_1 \to M \to C : (n, e) \mapsto m \mapsto m^e \bmod n$ where $M = (\mathbb{Z}/n\mathbb{Z})^\times = x : \mathbb{Z}/n\mathbb{Z} : \gcd(x, n) = 1 = C$
        * Weird that M depends on n (part of the key). In practice, it doesn't matter because the only messages that divide n are the primes, which breaks RSA anyways
    - $Dec : K_2 \to C \to M : (n, d) \mapsto c \mapsto c^d \bmod m$
- Correctness: Must show that $(m^e \bmod n)^d \bmod n = m$ *Proof.* $(m^e \bmod n)^d \bmod n = m^{ed} \bmod n$ (exponentiation under mod). Then, since $d = e^{-1} \bmod \phi(n)$, there exists k such that $de - 1 = k\phi(n)$, we have $m^{\phi(n)k+1} \equiv (m^{\phi(n)})^k m \equiv m \pmod{m}$. This holds by Euler's theorem ($\forall m \in (\mathbb{Z}/n\mathbb{Z})^\times, m^{\phi(n)} \equiv 1 \pmod{n}$) or Fermat's Little Theorem + Chinese Remainder Theorem (MATH 135)
- Security: Trivial that factoring $n = pq$ breaks RSA by computing $\phi(n)$
    - Conversely, if you know $\phi(n) = (p-1)(q-1)$ you can take $q\phi(n) = (n-1)(q-1)$ and solve for q
        * To avoid this, use the Carmichael exponent $\lambda(n) = \text{lcm}(p-1, q-1)$ instead of $\phi(n)$ which works. Of course, this doesn't work in practice because it's not actually that much different
    - For any non-trivial case, knowing one pair (e, d) also allows factoring n
    - Must make an assumption about hardness to prove security:
        * Factoring assumption: factoring random integers is hard
        * RSA factoring assumption: factoring $n = pq$ is hard (see, e.g., elliptical curve algorithm which depends on size of smallest prime in the factorization)
            · Of course, quantum computing fucks all of this to hell (see troll PQRSA which uses many small primes to make terabyte-sized moduli)
        * RSA assumption: given n, e, $m^e \bmod n$, it is hard to find m
    - Can prove RSA assumption $\implies$ RSA works (cannot prove without assumption without better results from complexity theory)

Lecture 4   Security Definitions (09/14)

- Security definitions, e.g., OW-CPA, IND-CPA, IND-CCA (Boneh, Shoup)
- How secure is a cryptosystem? Specify:
    - Allowable interactions between adversaries and parties
        * Second part of abbreviation
    - Computational limits of adversary
        * Not usually specified, usually probabilistic polynomial time

- Goal of the adversary to "break" the cryptosystem
  * First part of abbreviation
- OW-CPA: "one-way chosen-plaintext attack"
  - Adversary, given public key pk and encryption c of message m under pk, wants to determine m
  - Formally, given a random pk and c such that $c = Enc(pk, m)$ for some random m, it is infeasible for any probabilistic polynomial time algorithm $\mathcal{A}$ to determine m with non-negligible probability. That is, $\Pr[\mathcal{A}(pk, c) = m] = O(\frac{1}{\lambda^c})$ for all $c > 0$.
- Easier way to formalize ("Sequences of Games", Shoup 2004)
  - Two players: challenger $\mathcal{C}$ and adversary $\mathcal{A}$
  - Then, OW-CPA is
    1. $\mathcal{C}$ runs $KeyGen : \mathbb{1}^\lambda \xrightarrow{\$} (pk, sk)$
    2. $\mathcal{C}$ chooses $m \xleftarrow{\$} M$
    3. $\mathcal{C}$ computes $c \leftarrow Enc(pk, m)$
    4. $m' \xleftarrow{\$} \mathcal{A}(pk, c)$
    * with the win condition that $m' = m$, and we say that a cryptosystem is OW-CPA if a probabilistic polynomial time adversary $\mathcal{A}$ cannot win this game with non-negligible probability
  - IND-CPA (Goldmeier, Micoli 1984): indistinguishability
    1. $\mathcal{C}$ runs $(pk, sk) \xleftarrow{\$} KeyGen(\mathbb{1}^\lambda)$
    2. $(m_0, m_1) \xleftarrow{\$} \mathcal{A}(\mathbb{1}^\lambda, pk)$
    3. $\mathcal{C}$ picks $b \xleftarrow{\$} 0, 1$
    4. $\mathcal{C}$ computes $c \xleftarrow{\$} Enc(pk, m_b)$
    5. $b' \xleftarrow{\$} \mathcal{A}(\mathbb{1}^\lambda, pk, c)$
    * with the win condition $b = b'$, and a cryptosystem is IND-CPA if for all prob. poly. time $\mathcal{A}$, $\left|\frac{1}{2} - \Pr[\text{win}]\right| = O(\frac{1}{\lambda^\varepsilon})$ for all $\varepsilon > 0$
    * Encryption function must be random, otherwise $\mathcal{A}$ can re-encrypt

Lecture 5   Actual IND-CPA systems (09/16)

- IND-CPA is the standard security definition for symmetric security
  - Ciphertext contains no information about plaintext (except length)
- Design a slightly different equivalent IND-CPA game:
  1. $\mathcal{C}$ runs $(pk, sk) \xleftarrow{\$} KeyGen(\mathbb{1}^\lambda)$
  2. $(m_0, m_1) \xleftarrow{\$} \mathcal{A}(\mathbb{1}^\lambda, pk)$
  3. $\mathcal{C}$ picks $b \xleftarrow{\$} 0, 1$
  4. $\mathcal{C}$ computes $c_1 \xleftarrow{\$} Enc(pk, m_b)$ and $c_2 \xleftarrow{\$} Enc(pk, m_{b-1})$
  5. $b' \xleftarrow{\$} \mathcal{A}(\mathbb{1}^\lambda, pk, c_1, c_2)$
- Consider textbook RSA: $\mathcal{A}$ can choose $m_0 \neq m_1$ and compute $Enc(pk, m_0)$ and $Enc(pk, m_1)$ which allows it to win
  - In general, this applies to any scheme with deterministic encryption
- Goldwasser-Micali ("Probabilistic Encryption" 1982)
  1. Pick $n = pq$ (useful to have $p \equiv q \equiv 3 \pmod 4$)
  2. Pick $r \in (\mathbb{Z}/n\mathbb{Z})^\times$ such that $r \not\equiv x^2 \pmod p$ and $r \not\equiv x^2 \pmod q$
  3. Define $pk = (n, r)$ and $sk = (p, q)$
  4. Select a message bit b from $M = 0, 1$

5. Encrypt $Enc(b) = r^b y^2$ for some $y \overset{\$}{\leftarrow} (\mathbb{Z}/n\mathbb{Z})^\times$
   – Then, decrypt by determining ciphertext's squareness mod n
     * This is easy with the factorization $n = pq$ by Euler's criterion (a is square mod prime p if and only if $a^{(p-1)/2} \equiv 1 \pmod{p}$)
     * Determining squareness without factorization of n is hard, apparently
   – Since plaintexts are one bit, OW $\iff$ IND and this is provable under the circular-y assumption that determining squareness is hard
   – Also one bit messages are literally useless so who cares
- Elgamal (1984) (sometimes IND-CPA)
  – Publickeycryptosystemified Diffie-Hellman
  1. Setup is the same as DH, take some element $g \in G$ of a group
  2. Define $pk = g^x$ and $sk = x$
  3. Encrypt $Enc(m) = (g^y, g^{xy} \cdot m)$ for $y \overset{\$}{\leftarrow} \mathbb{Z}$
  – Then, decrypt $Dec(c_1, c_2) = \frac{c_2}{c_1^x} = \frac{g^{xy} \cdot m}{(g^y)^x} = m$
  – In general, key sharing schemes can be cryptosystemified like this
  – In an IND-CPA game, given $(g^y, g^{xy} m_b)$
    * Divide out m_0 to get either g^{xy} (if $m_b = m_0$) or garbage
    * Real challenge is distinguishing g^{xy} from garbage
  – Decisional Diffie-Hellman assumption: in the following game, $\left|\Pr[\mathcal{A} \text{ wins}] - \frac{1}{2}\right|$ is negligible in λ
    1. $\mathcal{C}$ chooses $p \overset{\$}{\leftarrow} \mathbb{Z}$ prime, $p \approx 2^\lambda$
    2. $\mathcal{C}$ chooses $g \in (\mathbb{Z}/p\mathbb{Z})^\times$
    3. $\mathcal{C}$ chooses $x, y \overset{\$}{\leftarrow} \mathbb{Z}$ and $h \overset{\$}{\leftarrow} (\mathbb{Z}/p\mathbb{Z})^\times$, computes $g_1 = g^x$, $g_2 = g^y$, $g_3 = g^{xy}$
    4. $\mathcal{C}$ chooses $b \overset{\$}{\leftarrow} 0, 1$ and $g_4 = g_3$ if $b = 0$ and h if $b = 1$
    5. $b' \leftarrow \mathcal{A}(\mathbb{1}^\lambda, p, g, g_1, g_2, g_4)$
  – Can prove: if DDH assumption holds, Elgamal is IND-CPA
- Layers of assumptions here:
  – DLOG: given g and g^x, it is hard to find x
  – CDH: given g, g^x, and g^y, it is hard to find g^{xy} (equivalent to Elgamal being OW-CPA)
  – DDH: given g^{xy} and garbage, is hard to distinguish the garbage
- How to piss off mathematicians: solving DLOG in $\mathbb{Z}/n\mathbb{Z}$ is easy but in $(\mathbb{Z}/p\mathbb{Z})^\times$ is hard
  – But $(\mathbb{Z}/p\mathbb{Z})^\times$ is isomorphic to $\mathbb{Z}/(p-1)\mathbb{Z}$ so DLOG difficulty must not be preserved over isomorphism
  – Specifically, DLOG is as exactly hard as computing the isomorphism (notice that we send $x \mapsto g^x$)
- DDH is actually easy in $(\mathbb{Z}/p\mathbb{Z})^\times$, need a subgroup $G \subset (\mathbb{Z}/p\mathbb{Z})^\times$ with $|G|$ prime

Chapter 2

Quadratic Residues

Lecture 6   Number Theory Background (09/19)

- Recall: RSA primes are gigantic so it takes time to do operations
  - e.g. picking $e \in (\mathbb{Z}/\phi(n)\mathbb{Z})^{\times}$ or finding $d = e^{-1} \pmod{\phi(n)}$ using EEA which runs in a logarithmic number of steps
  - e.g. running $Enc(m) = m^e \pmod{n}$ or $Dec(c) = c^d \pmod{n}$ using square-and-multiply which runs in a logarithmic number of steps
- Hard: picking non-squares in integers modulo p
  - Set of primes $\left|((\mathbb{Z}/p\mathbb{Z})^{\times})^2\right| = \frac{p-1}{2}$ for odd $p > 2$
  - This is because $f(x) = x^2$ is a 2-to-1 function on $(\mathbb{Z}/p\mathbb{Z})^{\times}$
    * To prove, show $f(a) = f(b) \iff a = \pm b$
    * Apply Euclid's Lemma: $p \mid (x - y)(x + y)$ implies $p \mid x - y$ or $p \mid x + y$, equivalently, $x = y \pmod{p}$ or $x = -y \pmod{p}$
    * Also another theorem: for R integral domain, every polynomial of degree n over R has at most n roots

Lecture 7   Squares Under a Modulus (09/21)

The big problem: Given $(\mathbb{Z}/n\mathbb{Z})^{\times}$ and $x \in (\mathbb{Z}/n\mathbb{Z})^{\times}$, when is $x \equiv \square \pmod{n}$?

For example, for $\mathbb{Z}/15\mathbb{Z}$, 1 and 4 are squares; for 8: just 1; for 7: 1, 2, and 4; and for 13: 1, 3, 4, 9, 10, and 12.

This breaks down into cases: n composite, n prime power, n prime

Theorem

> Suppose $n = \prod p_i^{e_i}$. Then, $x \equiv \square \pmod{n}$ if and only if for all i, $x \equiv \square \pmod{p_i^{e_i}}$.

Proof. Suppose $x = y^2 \pmod{n}$ for a unit y. Then, $n \mid (x - y^2)$ and $p_i^{e_i} \mid (x - y^2)$ by transitivity. That is, $x \equiv y^2 \pmod{p_i^{e_i}}$. In the reverse direction, if $p_i^{e_i} \mid (x - y^2)$ for all i, then by UPF (with some omitted detail), $n \mid (x - y^2)$. $\qquad\square$

The prime power case reduces to the prime case under condtions discovered in the homework problems lol.

Theorem

> The number of squares in $(\mathbb{Z}/p\mathbb{Z})^\times$ is $\frac{p-1}{2}$ for primes $p \geq 3$.

Proof. This is because $x = y^2 = (-y)^2$ and the size of the set is $p - 1$.

Build a table (x, g^x) instead of (x, x^2):

For $p = 13$ and $g = 2$, we get $(1, 2, 4, 8, 3, 6, 12 = -1, -2, -4, -8, -3, -6, -12 = 1)$ and the squares are the even-indexed values $(1, 4, 3, 12, 9, 10, 1)$.

This works for tables starting with non-squares: in fact, if $g \neq \square$, then $g^3 \neq \square$ (by the contrapositive, if $g^3 = \square$, then $g = \frac{g^3}{g^2} = \frac{\square}{\square} = \square$).

This gives us the result that $g^x = g^y$ when $x \equiv y \pmod{p-1}$ (note that this is equivalent to Fermat's Little Theorem, the reverse direction requires g coprime to $p-1$). $\qquad \square$

Definition (*order*)

> $\mathrm{ord}(a)$ is the period of $x \mapsto a^x$ for $a \in (\mathbb{Z}/p\mathbb{Z})^\times$.
> Equivalently, $\mathrm{ord}(a) = \min\{t \in \mathbb{Z} : a^t = 1, t > 0\}$.

Lemma

> Given elements a and b, numbers x and y:
> - $a^x = 1$ if and only if $\mathrm{ord}(a) \mid x$
> - $a^x = a^y$ if and only if $x \equiv y \pmod{\mathrm{ord}(a)}$
> - $\mathrm{ord}(a^x) = \frac{\mathrm{ord}(a)}{\gcd(x, \mathrm{ord}(a))}$
> - If $\mathrm{ord}(a)$ and $\mathrm{ord}(b)$ are coprime, then $\mathrm{ord}(ab) = \mathrm{ord}(a)\,\mathrm{ord}(b)$.

Proof. Only prove the last one:

Let $t = \mathrm{ord}(a)$, $u = \mathrm{ord}(b)$, $v = \mathrm{ord}(ab)$. Then, $(ab)^{tu} = a^{tu}b^{tu} = 1^u 1^t = 1$ so we have $v \mid tu$. Now, WLOG, $(ab)^{vu} = 1^u = 1 \implies a^{vu}b^{vu} = a^{vu}1 = a^{vu} = 1$. This gives $t \mid vu$ and $t \mid v$ since $\gcd(t, u) = 1$. Likewise, $u \mid v$ and we can conclude $tu \mid v$ because $\gcd(t, u) = 1$. That is, $tu = v$. $\qquad \square$

Lecture 8   Squares cont'd (09/23)

Definition (*primitive element*)

> $g \in G$ where $\{g^n : n \in \mathbb{N}\} = G$. Also called a generator.

Recall: if there exists primitive $g \in (\mathbb{Z}/p\mathbb{Z})^\times$, then for all $h \in (\mathbb{Z}/p\mathbb{Z})^\times$ where $h = g^k$, $h \equiv \square \iff k$ even. We can determine squareness using this fact, but finding k such that $h = g^k$ is doing a discrete log, which is hard.

Whether or not a primitive element exists is a non-trivial observation:

Theorem (*Gauss' primitive root*)

> For all primes p, $(\mathbb{Z}/p\mathbb{Z})^\times$ has a primitive element.

Proof. Observe that for all polynomials $f(x) \neq 0$ over $\mathbb{Z}/p\mathbb{Z}$, the number of roots of $f(x)$ is at most $\deg f$. Note that factorization fails in $\mathbb{Z}/n\mathbb{Z}$ in general: e.g. $x^2 - 1 = (x-1)(x+1) = (x-3)(x-5) \bmod 8$ or something weird like $x = (3x+2)(2x+3) \bmod 6$. We have this observation because $\mathbb{Z}/p\mathbb{Z}$ is an integral domain (and indeed, a field).

Consider $a \in (\mathbb{Z}/p\mathbb{Z})^\times$.

Claim $t = \operatorname{ord}(a) \mid p - 1$. Write $p - 1 = tq + r$. If $r = 0$, done. If $r > 0$, $\operatorname{ord}(a) = r < t$, contradiction and indeed $r = 0$.

For each divisor d of $p-1$, consider $S_d = \{x \in (\mathbb{Z}/p\mathbb{Z})^\times : \operatorname{ord}(x) = d\}$. Then, $\bigcup_{d|p-1} S_d = (\mathbb{Z}/p\mathbb{Z})^\times$ and this is a disjoint union. To prove Gauss' theorem, we just need $\left|S_{p-1}\right| > 0$.

Proceed in general for arbitrary $|S_d| > 0$ for all $d \mid p - 1$.

If $S_d = \varnothing$, then $|S_d| = 0$. Otherwise, claim that $|S_d| = \phi(d) = |(\mathbb{Z}/d\mathbb{Z})^\times|$.

If S_d is not empty, then $\exists a \in S_d$ where $\operatorname{ord}(a) = d$. Consider $x^d - 1$. The roots of this polynomial will include all elements of S_d (and others). We can write the set of roots as exactly $\{a^0, \dots, a^{d-1}\}$. So for all $b \in S_d$, $b = a^k$ since b is a root and we need only count those powers with order d. But that is exactly $\operatorname{ord}(a^i) = \frac{\operatorname{ord}(a)}{\gcd(i,d)} = \frac{d}{\gcd(i,d)}$. So we are counting the i such that $\gcd(i, d) = 1$, which is exactly $\phi(d)$.

Now, $p - 1 = |(\mathbb{Z}/p\mathbb{Z})^\times| = \left|\bigcup_{d|p-1} S_d\right| = \sum |S_d| \leq \sum \phi(d)$ which is equal to $p - 1$ by Möbius inversion. That last inequality being an equality implies that $|S_d| \neq 0$ for any $d \mid p - 1$, and in particular $p - 1 \mid p - 1$.

Quick combinatorical proof of this fact: write out all the $p - 1$ fractions over $p - 1$, then each of $\phi(d)$ is the number of fractions where the denominator reduces to d. The sum must be $p - 1$. $\qquad \square$

Lecture 9   Applying to DDH (09/26)

Recall the Decisional Diffie-Hellman problem: Given g, g^x, g^y, g^z, determine if $z = xy$. Formally, as a game:

- $\mathcal{C}$ chooses a bit $b \in \{0, 1\}$ and $x, y \xleftarrow{\$} \mathbb{Z}$
- $b' \leftarrow \mathcal{A}(g, g^x, g^y, g^z)$ where $z \leftarrow \begin{cases} xy & b = 0 \\ \xleftarrow{\$} \mathbb{Z} & b = 1 \end{cases}$
- Win condition: $b = b'$ with non-negligible probability

Notice that if g is a primitive root, then $|\{g^x : x \in \mathbb{Z}\}| = p - 1$. But bruteforce DLOG takes $\frac{p-1}{2}$ steps on average. Then, Elgamal is IND-CPA $\iff$ DDH holds.

Proposition

The Decisional Diffie-Hellman assumption in $(\mathbb{Z}/p\mathbb{Z})^\times$ with a primitive base g does not hold.

Proof. We tell squares and non-squares apart.

Recall from last lecture's theorem we have that if g is a primitive root, $g^x \equiv \square \pmod{p} \iff x \equiv 0 \pmod{2}$. Then, by Euler's criterion, $a \equiv \square \pmod{p} \iff a^{(p-1)/2} \equiv 1 \pmod{p}$. Therefore, it is possible to tell the parity of x, y, and z in reasonable time using Euler's criterion (since raising to a power is easy).

If xy is odd only when x and y are odd, so if you know the parity of z you can distinguish if $z = xy$ or random with non-negligible advantage. $\qquad \square$

Proposition (*Euler's criterion*)

$$a \equiv \square \pmod{p} \iff a^{(p-1)/2} \equiv 1 \pmod{p}$$

Proof. Suppose $a \equiv \square$ iff $a \equiv g^k$ for even $k = 2\ell$ iff $a^{(p-1)/2} = (g^k)^{(p-1)/2} = g^{k(p-1)/2} = (g^{p-1})^\ell = 1^\ell = 1$ by FℓT.

Otherwise, $a \not\equiv \square$ iff $a = g^k$ for $k = 2\ell + 1$ iff $a^{(p-1)/2} = (g^k)^{(p-1)/2} = g^{(p-1)/2\cdot(2\ell+1)} = g^{(p-1)/2\cdot2\ell} \cdot g^{(p-1)/2} = g^{(p-1)/2} \neq 1$. But in fact $g^{(p-1)/2} = \sqrt{g^{p-1}} = \sqrt{1} = -1$ since it is not positive 1. $\square$

Corollary. For $p > 2$, -1 is a square mod p if and only if $p \equiv 1 \pmod{4}$.

Proof. For -1 to be a square, we need $(-1)^{(p-1)/2} \equiv 1 \pmod{p}$. That is, $\frac{p-1}{2}$ is even and we have $p \equiv 1 \pmod{4}$. $\square$

This quantity $g^{(p-1)/2}$ is useful and we give it a name:

Definition (*Legendre symbol*)

For $p > 2$ and $a \in \mathbb{Z}/p\mathbb{Z}$, the quadratic character of a, written $(\frac{a}{p}) = a^{(p-1)/2}$, is 1 if $a \equiv \square$, 0 if $a \equiv 0$, and -1 if $a \not\equiv \square$.

Equivalently, define $\chi_p : (\mathbb{Z}/p\mathbb{Z})^\times \to \{\pm1\} : a \mapsto (\frac{a}{p})$ and notice that this is a multiplicative homomorphism that preserves $\chi_p(ab) = \chi_p(a)\chi_p(b)$.

Theorem (*multiplicativity*)

$$(\tfrac{ab}{p}) = (\tfrac{a}{p})(\tfrac{b}{p})$$

Proof. $(\frac{ab}{p}) = (ab)^{(p-1)/2} = a^{(p-1)/2}b^{(p-1)/2} = (\frac{a}{p})(\frac{b}{p})$ $\square$

Lecture 10   Quadratic Characters in the Complex Plane (09/28)

Recall: we have that for odd primes, $(\frac{-1}{p}) = 1 \iff p \equiv 1 \pmod{4}$ which we proved by applying Euler's criteron. We have the similar lemma:

Lemma

$$(\tfrac{2}{p}) = 1 \iff p \equiv 1, 7 \pmod{8}.$$

Proof. This is harder because $2^{(p-1)/2}$ is not easy to analyze, i.e., the order of 2 is not easy to derive.

What numbers, in general, have finite/known order? Complex roots of unity $\zeta_n = e^{2\pi i/n}$.

We can write $\sqrt{2} = \zeta_8 + \zeta_8^7$, so $2^{(p-1)/2} = (\zeta_8 + \zeta_8^7)^{p-1} = \frac{(\zeta_8+\zeta_8^7)^p}{\zeta_8+\zeta_8^7}$. The last transformation is helpful since p powers behave well mod p.

Now, notice that $(x+y)^p \equiv x^p + y^p \pmod{p}$ because all the other terms will have a factor of $p \mid \binom{p}{i}$.

Therefore, $(\frac{2}{p}) \equiv 2^{(p-1)/2} \equiv \frac{\zeta_8^p+\zeta_8^{7p}}{\zeta_8+\zeta_8^7} \pmod{p}$.

There are four cases for $p \pmod 8$ because we assume $p > 2$:

1. $\zeta_8^p = \zeta_8^1$ and $\zeta_8^{7p} = \zeta_8^7$

3. $\zeta_8^p = \zeta_8^3$ and $\zeta_8^{7p} = \zeta_8^5$
5. $\zeta_8^p = \zeta_8^5$ and $\zeta_8^{7p} = \zeta_8^3$
7. $\zeta_8^p = \zeta_8^7$ and $\zeta_8^{7p} = \zeta_8^1$

Clearly, for $p \equiv 1, 7 \pmod 8$, we have $\frac{\zeta_8^p + \zeta_8^{7p}}{\zeta_8 + \zeta_8^7} = \frac{\zeta_8 + \zeta_8^7}{\zeta_8 + \zeta_8^7} = 1$. Slightly less intuitively, for $p \equiv 3, 5 \pmod 8$, notice that $\zeta_8^3 + \zeta_8^5 = -\sqrt{2}$, so the fractions go to -1. $\square$

Note: We can algebraically extend $\mathbb{Z}/p\mathbb{Z}$ with the necessary complex numbers to make the proof valid (or simply assert that the necessary roots of unity exist).

The pattern sort of extends:

- $(\frac{3}{p}) = 1$ if $p = 1, 11$ mod 12 and -1 if $p = 5, 7$ mod 12.
- $(\frac{5}{p}) = 1$ if $p = \pm 1, \pm 9$ mod 20 and -1 if $p = \pm 3, \pm 7$ mod 20.
- $(\frac{7}{p}) = 1$ if $p = \pm 1, \pm 3, \pm 9$ mod 28 and -1 if $p = \pm 5, \pm 11, \pm 13$ mod 28.

In fact, we have $(\frac{7}{p}) = 1$ if $p = \pm 1, \pm 9, \pm 25$ mod 28. This flips the question from is 7 a square mod p to asking if p is a square mod 28.

Lemma

$$(\frac{-3}{p}) = \begin{cases} 1 & p \equiv 1 \pmod 3 \\ -1 & p \equiv 2 \pmod 3 \end{cases}$$

Proof. Consider again $(-3)^{(p-1)/2} = (\sqrt{-3})^{p-1}$. We can notice $\sqrt{-3} = \sqrt{3}i = \zeta_6 + \zeta_3$.

This gives us $(\sqrt{-3})^{p-1} = \frac{\zeta_3^p - \zeta_3^{2p}}{\zeta_3 - \zeta_3^2}$ because $\zeta_6 = -\zeta_3^2$.

If $p \equiv 1 \pmod 3$, then $\frac{\zeta_3^p - \zeta_3^{2p}}{\zeta_3 - \zeta_3^2} = \frac{\zeta_3 - \zeta_3^2}{\zeta_3 - \zeta_3^2} = 1$ and if $p \equiv 2 \pmod 3$, $\frac{\zeta_3^p - \zeta_3^{2p}}{\zeta_3 - \zeta_3^2} = \frac{\zeta_3^2 - \zeta_3^1}{\zeta_3 - \zeta_3^2} = -1$. $\square$

Notice that to get to $\sqrt{3}$ on the complex plane, we need ζ_{12}, which explains why we see mod 12 in the rule. To get $\sqrt{5}$, we can either use the fact that $\cos \frac{2\pi}{5} = \frac{1}{4}(\sqrt{5} - 1)$ or notice that $(\zeta_5 - \zeta_5^2 - \zeta_5^3 + \zeta_5^4)^2 = (4 - \zeta_5 - \zeta_5^2 - \zeta_5^3 - \zeta_5^4) = 5 - (1 + \zeta_5^1 + \zeta_5^2 + \zeta_5^3 + \zeta_5^4) = 5$. We can then execute the same fraction-by-cases technique, getting our result mod 5.

Aside: This is the Gauss sum for $\sqrt{5} = \sum (\frac{i}{5}) \zeta_5^i$.

Lecture 11   Quadratic Reciprocity (09/30)

Recall the pattern from last lecture, where we noticed that asking if q is a square mod p seems to be like asking if p is a square mod $4q$. This is almost true, but in fact

Theorem (*Quadratic Reciprocity*)

$(\frac{q}{p}) = (\frac{p}{q})$ for odd primes $p \neq \pm q$ where at least one is congruent to 1 mod 4 and at least one is positive.

Equivalently, for all distinct positive odd primes p and q, $(\frac{p}{q})(\frac{q}{p}) = (-1)^{\frac{p-1}{2}\frac{q-1}{2}}$

The proof follows by Gauss sums and the vague ideas from the last lecture.

This means we can evaluate any Legendre symbol using a modulus as either one of

$$\left(\frac{-1}{p}\right) = (-1)^{\frac{p-1}{2}} = \begin{cases} 1 & p \equiv 1 \pmod 4 \\ -1 & p \equiv 3 \pmod 4 \end{cases}$$

$$\left(\frac{2}{p}\right) = (-1)^{\frac{p^2-1}{8}} = \begin{cases} 1 & p \equiv \pm 1 \pmod 8 \\ -1 & p \equiv \pm 3 \pmod 8 \end{cases}$$

$$\left(\frac{q}{p}\right) = \left(\frac{p}{q}\right)(-1)^{\frac{p-1}{2}\frac{q-1}{2}} = \begin{cases} (\frac{p}{q}) & p \equiv \pm 1 \pmod 4 \text{ or } q \equiv \pm 1 \pmod 4 \\ -(\frac{p}{q}) & p \equiv q \equiv 3 \pmod 4 \end{cases}$$

which is nicer than using Euler's criterion.

Example 11.1. Is 71 a square mod 101?

Solution. Write $\left(\frac{71}{101}\right) = \left(\frac{101}{71}\right) = \left(\frac{30}{71}\right) = \left(\frac{2}{71}\right)\left(\frac{3}{71}\right)\left(\frac{5}{71}\right)$ by quadratic reciprocity and multiplicativity.

Then, $\left(\frac{2}{71}\right) = 1$ since $71 \equiv 7 \pmod 8$.

Also, $\left(\frac{3}{71}\right) = -\left(\frac{71}{3}\right) = -\left(\frac{2}{3}\right) = 1$ since $71 \equiv 3 \pmod 4$.

Finally, $\left(\frac{5}{71}\right) = \left(\frac{71}{5}\right) = \left(\frac{1}{5}\right) = 1$ since 1 is always a square.

This gives $\left(\frac{71}{101}\right) = 1 \cdot 1 \cdot 1 = 1$ so 71 is a square mod 101. $\qquad\square$

Asymptotically, this is not faster than Euler's criterion because we require factoring. However, it is prettier.

To deal with a random large number, we must consider what to do after factoring out all the 2s (since we can deal with those quickly).

Definition (*Jacobi symbol*)

> For all $m, n \in \mathbb{N}_{>0}$ with n odd, $\left(\frac{m}{n}\right) = \prod_{i=1}^{k}\left(\frac{m}{p_i}\right)$ where $\prod_{i=1}^{k} p_i = n$ is the prime factorization of n

Theorem (*Jacobi*)

> For all positive and odd m and n,
>
> $$\left(\frac{-1}{n}\right) = (-1)^{\frac{n-1}{2}} = \begin{cases} 1 & n \equiv 1 \pmod 4 \\ -1 & n \equiv 3 \pmod 4 \end{cases}$$
>
> $$\left(\frac{2}{n}\right) = (-1)^{\frac{n^2-1}{8}} = \begin{cases} 1 & n \equiv \pm 1 \pmod 8 \\ -1 & n \equiv \pm 3 \pmod 8 \end{cases}$$
>
> $$\left(\frac{m}{n}\right) = \left(\frac{n}{m}\right)(-1)^{\frac{n-1}{2}\frac{m-1}{2}} = \begin{cases} (\frac{n}{m}) & n \equiv \pm 1 \pmod 4 \text{ or } m \equiv \pm 1 \pmod 4 \\ 0 & \gcd(m,n) \neq 1 \\ -(\frac{n}{m}) & n \equiv m \equiv 3 \pmod 4 \end{cases}$$

Note: For Legendre symbols, $\left(\frac{a}{p}\right) = 1 \iff a \equiv \square \pmod p$. However, for Jacobi symbols, we only have the one-way implication $\left(\frac{m}{n}\right) = -1 \implies m \not\equiv \square \pmod n$.

Return now to the application to cryptography, specifically to Goldwasser–Micali.

Goldwasser–Micali cryptosystem

Key Generation: Choose random primes p, q. Set $n = pq$.

Choose $x \in (\mathbb{Z}/n\mathbb{Z})^\times$ such that $\left(\frac{x}{p}\right) = \left(\frac{x}{q}\right) = -1$ (then $\left(\frac{x}{n}\right) = 1$). Publish x.

Encrypt: $m \in \{0, 1\}$

Choose some $r \xleftarrow{\$} (\mathbb{Z}/n\mathbb{Z})^\times$. Then, $Enc(m) = x^m r^2 = c$.

Decrypt: Determine whether c is a "fake" square using the factorization.

The underlying assumption is that it is not easy to distinguish actual squares mod n and "fake" squares mod n.

Chapter 3

Primality

Lecture 12   Primality Testing (10/03)

Given $n \in \mathbb{Z}$, how can we tell if n is prime?

Lemma (*Fermat test*)

> Recall FℓT: for a prime p, $a \in (\mathbb{Z}/p\mathbb{Z})^\times \implies a^{p-1} = 1$. Therefore, if $a \in (\mathbb{Z}/n\mathbb{Z})^\times$ and $a^{n-1} \neq 1$, then n is not prime.

Definition (*Fermat witness*)

> Let $n \in \mathbb{N}$, $\alpha \in (\mathbb{Z}/n\mathbb{Z})^\times$ where $\alpha^{n-1} \neq 1$.

When n is prime, no Fermat witness can exist. When n is not prime, only some elements are Fermat witnesses. The other elements are *Fermat liars*. How many liars are in $(\mathbb{Z}/n\mathbb{Z})^\times$?

Theorem

> For $n > 2$, if there exists one Fermat witness in $(\mathbb{Z}/n\mathbb{Z})^\times$, then there exist at least $\frac{\phi(n)}{2}$ Fermat witnesses.

Proof. Consider the set $H = \{\alpha \in (\mathbb{Z}/n\mathbb{Z})^\times : \alpha^{n-1} = 1\}$.

H is a subgroup: $1 \in H$, $ab \in H$, $a^{-1} \in H$ (trivial by exponentiation properties).

So by Lagrange's theorem, $|H| \mid |(\mathbb{Z}/n\mathbb{Z})^\times|$.

Either (1) $|H| = \phi(n)$, so there are no witnesses, or (2) $|H| < \phi(n)$, so $|H| \leq \frac{\phi(n)}{2}$. $\qquad\square$

Definition (*Carmichael number*)

> $n \in \mathbb{N}$, $n > 2$ such that n is composite and n has no Fermat witnesses.

Examples: $n = 561 = 3 \times 11 \times 17$. By FℓT, we have $\alpha^{n-1} = \alpha^{560}$ is 1 mod 3, 1 mod 11, and 1 mod 17.

Recall that for n prime: $a^{\frac{n-1}{2}} \equiv \left(\frac{a}{n}\right) \pmod{n}$ when $n > 2$, odd, and $a \in (\mathbb{Z}/n\mathbb{Z})^\times$. This gives us the following test:

Lemma (*Solovay–Strassen test*)

> If $a^{\frac{n-1}{2}} \not\equiv \left(\frac{a}{n}\right) \pmod{n}$, then n is not prime.

We can calculate $a^{\frac{n-1}{2}}$ by repeated squaring and $\left(\frac{a}{n}\right)$ by Jacobi reciprocity and factoring out 2's. We can now define witnesses as in the Fermat test.

Definition (*Euler (Solovay–Strassen) witness*)

> An element $\alpha \in (\mathbb{Z}/n\mathbb{Z})^\times$ where $\left(\frac{\alpha}{n}\right) \not\equiv \alpha^{\frac{n-1}{2}} \pmod{n}$. If an element is not an Euler witness, it is an Euler liar.

Notice that all Euler witnesses must also be Fermat witnesses, meaning that hopefully we have a more refined test here.

Theorem

> If $n > 2$ is composite and odd, then there exists at least one Euler witness.

Proof. Suppose n is composite and $n = p \times k$.

If $p \nmid k$, then solve $\alpha \equiv \beta \pmod{p}$ and $\alpha \equiv 1 \pmod{k}$ where β is a quadratic non-residue mod p. Now, calculate

$$\left(\frac{\alpha}{n}\right) = \left(\frac{\alpha}{p}\right)\left(\frac{\alpha}{k}\right) = \left(\frac{\beta}{p}\right)\left(\frac{1}{k}\right) = (-1)(1) = -1$$

Suppose $\alpha^{\frac{n-1}{2}}$ is -1. Then, $\alpha^{\frac{n-1}{2}} \equiv -1 \pmod{n}$ and that means $\alpha^{\frac{n-1}{2}} \equiv -1 \pmod{k}$. But we know $\alpha \equiv 1 \pmod{k}$, so this is a contradiction.

Otherwise, $p \mid k$. Let $\alpha = 1 + k$. Calculate

$$\left(\frac{\alpha}{n}\right) = \left(\frac{1+k}{n}\right) = \left(\frac{1+k}{p}\right)\left(\frac{1+k}{k}\right) = \left(\frac{1}{p}\right)\left(\frac{1}{k}\right) = (1)(1) = 1$$

Suppose $\alpha^{\frac{n-1}{2}} = 1$. This implies that $\operatorname{ord}(\alpha) \mid \frac{n-1}{2}$. Calculate $\alpha^p = (1+k)^p = 1^p + \underbrace{pk^1 + \cdots + \binom{p}{p}k^p}_{0 \pmod{n}} = 1$ which implies $\operatorname{ord}(\alpha) = p$. But $p \mid n \implies p \nmid n - 1 \implies p \nmid \frac{n-1}{2}$.

Therefore, α is an Euler witness. $\qquad\square$

This theorem combined with the at-least-$\frac{\phi(n)}{2}$ theorem means that we have for every odd, composite $n > 2$ there are $\frac{\phi(n)}{2}$ Euler witnesses.

Lecture 13   Strong Primality Testing (10/05)

Recall: for the Fermat test, evaluate a^{n-1} a bunch of times. If it is equal to 1, prime or liar; otherwise, composite. For the Solovay–Strassen test, evaluate $a^{\frac{n-1}{2}} = \left(\frac{a}{n}\right)$. If yes, prime or Euler liar; otherwise, composite. Also, there are an infinite number of Carmichael numbers that screw with this but otherwise you have around a 50% chance of getting a witness.

We can refine this further beyond considering $n - 1$ and $\frac{n-1}{2}$.

Write $n - 1 = 2^t \cdot s$ so that s is odd. Then, a^{n-1} is a^s squared t times. So instead of asking if $a^{2^t s} = 1$, consider if $a^{2^{t-1}s}$ is an "expected" square root of 1, i.e., ± 1. If it is not,

it is composite. If it is and it is -1, we have a prime or liar. If it is and it is 1, keep going back. If we reach $a^s = 1$, we get no information.

Lemma (*Miller–Rabin test*)

> Let $x \leftarrow a^s$. Do:
> - If $x = 1$, stop. Probably prime.
> - If $x = -1$, stop. Probably prime.
> - Otherwise, $x \leftarrow x^2$
>
> while $x \neq a^{2^t s}$. If we reach the end, it is composite.

Definition (*Miller–Rabin (strong) liar*)

> $a \in (\mathbb{Z}/n\mathbb{Z})^\times$ if either $a^s = 1$ or $a^{2^k s} = -1$ for $0 \leq k < t$.

We call this a "strong liar" because every strong liar is an Euler liar, and every Euler liar is a Fermat liar.

Theorem

> Suppose n has at least two distinct prime factors. Then, the number of Miller–Rabin liars is at most $\frac{\phi(n)}{4}$ and in general, if n has ℓ distinct prime factors, there are at most $\frac{\phi(n)}{2^\ell}$ Miller–Rabin liars.

We can make these primality tests deterministic by iterating $a = 1, \dots, n$. We do not need to go to $a = n$ and instead we can establish an upper bound on the smallest witness. The bound (by Bach) is $O(\log^2 n)$, specifically, $2 \log^2 n$. But this requires the Generalized Riemann Hypothesis which everyone believes anyways, so we just check $a = 1, \dots, 2 \log^2 n$.

To analyze complexity, notice that we have $\log n$ multiplications at each step, i.e., $\log^{1+\epsilon} n$ bit operations using fast multiplication. So the complexity is $O(\log^{2+(1+\epsilon)+1} n)$.

Further reading:

- AKS (Agrawal–Kayal–Saxena; 2004) primality test in $O(\log^6 n)$ which does not rely on GRH and was an undegrad project(!!)
- ECPP (elliptic curve prime proving) notable for not having liars, also does not require GRH and runs non-deterministically (Monte Carlo) in $O(\log^5 n)$
- Cyclotomic primality test in $O((\log n)^{\log \log n})$, best until AKS proved that primality is in P.

Since there are $\frac{n}{\log n} + O(\sqrt{n})$ primes less than n, we can pick random numbers of size e^ℓ to get an approximate $\frac{1}{\ell}$ probability of a prime.

Lecture 14   Malleability (10/07)

Recall the Goldwasser–Micali cryptosystem. It satisfies IND-CPA provided that the quadratic reciprocity problem is hard. That is, determining whether an $x = pq$ with $\left(\frac{x}{n}\right) = 1$ is actually a square or not (i.e. $\left(\frac{x}{p}\right) = 1$).

However, an adversary can still alter the message without needing to decrypt. This also applies, for example, to XOR one-time pads (since if $c = k \oplus m$ and we intercept $c \mapsto c \oplus n$, recepient will get $m' = m \oplus n$). Using MACs can get around this problem (e.g. AES with GCM or Chacha20 with Poly1305).

Definition (*non-malleability*)

> Given the game NM-CPA:
> 1. $\mathcal{C}$ generates (pk, sk)
> 2. $(m_0, m_1, m_0', m_1') \xleftarrow{\$} \mathcal{A}(\lambda, pk)$ where $m_0' \neq m_1'$
> 3. $\mathcal{C}$ chooses $b \xleftarrow{\$} \{0, 1\}$
> 4. $\mathcal{C}$ computes $c = E(m_b)$
> 5. $c' \leftarrow \mathcal{A}(\lambda, pk, c)$
>
> with win condition $D(c') = m_b'$ with non-negligible probability above 50%.

Instead of CPA games, consider CCA2 (chosen-ciphertext attack 2) games. Here, the adversary has a decryption oracle that takes anything except c. In CCA1, the oracle can only be accessed prior to receiving c.

Theorem

> IND-CCA2 is equivalent to NM-CCA2

Note that IND-CPA is not equivalent to NM-CPA, which is instead equal to IND-PCA (parallel ciphertext attack, where all oracle queries must occur at once).

Lecture 15   Factorization Algorithms (10/17)

Naive approach: trial division by $1, \ldots, \sqrt{n}$ which is $O(\sqrt{n}) = O(\exp(\frac{1}{2} \log n))$. Note that we call this "exponential" because we measure with respect to the size of the input, i.e., $\lg n \approx \log n$.

Proposition

> If $x, y \in (\mathbb{Z}/n\mathbb{Z})^\times$ satisfy $x^2 \equiv y^2 \pmod{n}$ and $x \not\equiv \pm y \pmod{n}$, then $\gcd(n, x-y)$ is a non-trivial factor of n.

Proof. Since $x^2 - y^2 \equiv 0$, we have $(x - y)(x + y) \equiv 0$. But we know that $x - y \not\equiv 0$ and $x + y \not\equiv 0$ so there must be some weird hidden factor.

If $\gcd(n, x - y) = n$, then $n \mid (x - y) \implies x \equiv y \pmod{n}$ and if $\gcd(n, x - y) = 1$, then $n \mid (x - y)(x + y)$ which implies $n \mid (x + y)$ by Gauss' Lemma which gives the same contradiction. Therefore, since the GCD must divide n, it is non-trivial. $\qquad\square$

Using this, we can find non-trivial factors of n by finding x and y and then applying the EEA. How to find x and y?

Random squares (Dixon)

Definition (*B-smoothness*)

> $n \in \mathbb{N}$ where the largest prime factor is less than B

Chose $x_i \xleftarrow{\$} (\mathbb{Z}/n\mathbb{Z})^\times$. For each x_i, compute $x_i^2 \pmod{n}$ and keep the B-smooth squares. We can tell if a number is B-smooth by trial division (since B is small).

We need at least $t + 1$ squares that are B-smooth.[1]

---

[1]Where $t = \pi(B)$ is the prime-counting function.

This gives us squares $x_1^2 \bmod n = p_1^{e_{1,1}} \cdots p_t^{e_{t,1}}$ up to $x_{t+1}^2 \bmod n = p_1^{e_{1,t+1}} \cdots p_t^{e_{t,t+1}}$.

Take the subset product $\prod x_i^2 \bmod n = p_1^{\sum e_{1,i}} \cdots p_t^{\sum e_{t,i}}$.

We can define $b_i = \begin{cases} 0 & i \in S \\ 1 & i \notin S \end{cases}$ so that $\sum_{i \in S} e_{j,i} = \sum_{i=1}^{t+1} e_{j,i} b_i = 0 \bmod 2$ to find squares.

Solve this homogeneous linear system over $\mathbb{Z}/2\mathbb{Z}$ (where the b_i are variables). We know there exists a non-trivial solution because there are more variables (at least $t + 1$) than equations (exactly t).

That gives a square subset product $x^2 = \prod_{i=1}^{t+1} (x_i^2)^{b_i} \bmod n = \prod_{j=1}^{t} p_j^{\sum_{i=1}^{t+1} e_{j,i} b_i} \bmod n = y^2$.

The LHS and RHS are unrelated except for the fact that they are equal mod n. In fact, with about 50% probability, $x \not\equiv \pm y \bmod n$. The probability can be improved by increasing $t + 1$ to like $t + 10$. Since $t \approx B$ is large, this is negligible.

Picking B: large B makes it more likely to find B-smooth squares, however, the amount of work $t + 1$ is proportional to B.

We want to pick B such that the probability of squares being B-smooth is $\frac{1}{B}$. This depends on n.

From analytic number theory, the probability that a random $y \overset{\$}{\leftarrow} \mathbb{Z}/n\mathbb{Z}$ is $L(\alpha, c)$-smooth is $L(1 - \alpha, \frac{1-\alpha}{c})$.[2] So we set a bound on B of $L_n(\frac{1}{2}, \frac{\sqrt{2}}{2})$. Since $(\log n)^k \ll B \ll \sqrt{n}$, we call this subexponential.

Lecture 16   Better Sieves (10/19)

What is the probability that a particular $x^2 \bmod n$ is B-smooth? Vanishingly small for large n (in the hundreds of digits) and small-ish B (around 10^9). However, we can prove that the runtime for random squares is $L_n(\frac{1}{2}, 2\sqrt{2})$ using results from analytic number theory, i.e., probabilistic subexponential time.

How can we improve? Pick x such that $x^2 \bmod n$ is small (and more likely to be B-smooth). Naively: small numbers stay small (but are useless). Instead, pick $x \approx \sqrt{n}$ so that $x^2 \bmod n = x^2 - n$.

Then, if $x = \sqrt{n} + k$, $x^2 = (\sqrt{n} + k)^2 - n = 2k\sqrt{n} + k^2 = O(k\sqrt{n})$, i.e., around half the size of n and much smaller than n.

This is the quadratic sieve. We can bound $B < L_n(\frac{1}{2}, 1)$ and prove runtime $L_n(\frac{1}{2}, \sqrt{2})$.

Suppose we write $\sqrt{n} = a_0 + \cfrac{1}{a_1 + \cfrac{1}{a_2 + \frac{1}{\ddots}}} = [a_0, a_1, \dots]$ as a continued fraction.

Define $\frac{P_i}{Q_i} = [a_0, \dots, a_i]$. These fractions rapidly approach $\sqrt{n}$ (and are in fact the best rational approximations). That is, $P_i^2 - nQ_i^2$ rapidly approaches 0. We can prove that $0 < P_i - nQ_i^2 < 2\sqrt{n} + 1$. Then, we can take $P_i^2 \bmod n$ and sieve guaranteed that the squares are $O(2\sqrt{n})$.

---

[2]Where $L_n(\alpha, c) = O(\exp(c(\log n)^\alpha (\log \log n)^{1-\alpha}))$. Notably, $L_n(1, 1) = n$ and $L_n(1, c) = n^c$. Then, $\sqrt{n} = L_n(1, \frac{1}{2})$. Also, $L_n(0, c) = (\log n)^c$. That is, we interpolate between $L_n(0, c)$ polynomial time and $L_n(1, c)$ exponential time

Comparing the continued fraction sieve and quadratic sieve, $O(2\sqrt{n})$ appears better than $O(k\sqrt{n})$. However, if the quadratic sieve considers consecutive numbers to square, we can do a sieve of Eratosthenes-like search to find good B-smooth candidates. This is faster.

One more improvement step: number field sieve.

Choose $d \approx 6 \in \mathbb{Z}$ and $m \approx n^{1/d}$. Write n in base m: $n = a_0 + a_1 m + \cdots + a_5 m^5$ and consider the polynomials $f(x) = a_0 + a_1 x + \cdots + a_5 x^5$ and $g(x) = x - m$.

We know that $f(m) = n \equiv 0 \pmod{n}$ and $g(m) = 0$. That is, m is a root of both f and g mod n. The coefficients are also all around $m = O(n^{1/d})$ in size.

Consider α a complex root of f (but consider it as part of a number field $\alpha \in \mathbb{Z}[\alpha]$). We pick $a_i, b_i \in \mathbb{Z}$ such that $a_i + b_i \alpha = \prod \beta_i$ is smooth in $\mathbb{Z}[\alpha]$ and $a_i + b_i m = \prod q_i$ is smooth in $\mathbb{Z}$.

Pick a subset S such that $\prod_{i \in S}(a_i + b_i \alpha) = \square$ in $\mathbb{Z}[\alpha]$ and $\prod_{i \in S}(a_i + b_i m) = \square \bmod n$ in $\mathbb{Z}$. We can expand the first sum, then replace α with m mod n to get congruent squares for a sieve. In fact, $\alpha \mapsto m \bmod n$ is a ring homomorphism from $\mathbb{Z}[\alpha] \to \mathbb{Z}/n\mathbb{Z}$.

Since the numbers are smaller, we have complexity $L_n(\frac{1}{3}, \sqrt[3]{\frac{64}{9}})$.

Lecture 17   (10/21)

Lecture 18   Index Calculus (10/26)

There is a connection between runtimes of factoring algorithms and DLOG algorithms:

Factoring n	DLOG in $\mathbb{Z}_p^*$ or $\mathbb{F}_q^*$ where $q = p^k$
Trial ($O(n)$)	Trial ($O(p)$)
$O(\sqrt{n})$	Pollard's rho ($O(\sqrt{p})$)
Random squares ($L_p(\frac{1}{2}, \sqrt{2})$)	Index calculus ($L_p(\frac{1}{2}, \sqrt{2})$)
NFS ($L_n(\frac{1}{3}, \sqrt[3]{\frac{64}{9}})$)	NFS for DLOG ($L_p(\frac{1}{3}, \sqrt[3]{\frac{64}{9}})$)
Special NFS ($L_n(\frac{1}{3}, \sqrt[3]{\frac{32}{9}})$)	Tower NFS ($L_q(\frac{1}{3}, \sqrt[3]{\frac{32}{9}})$ if $k < 50$)
???	$L_q(\varepsilon, c)$ for $p < 10$

Theorem (*Shoup*)

> For a generic group, classical probabilistic DLOG algorithms require $\Omega(\sqrt{p})$ group operations.

What we mean by generic here is that the group "interface" is exposed (multiplication, inversion, equality) but we don't know anything about the elements/structure.

Index calculus

Consider $\mathbb{Z}_p^* = \{1, 2, \ldots, p-1\}$, g, $h = g^\alpha$. We want to find α, the "index". We construct "random index calculus" from the random squares algorithm. Pick random x_i and

calculate:

$$g^{x_1} \bmod p = p_1^{e_{1,1}} \dots p_t^{e_{t,1}}$$
$$\vdots$$
$$g^{x_{t+1}} \bmod p = p_1^{e_{1,t+1}} \dots p_t^{e_{t,t+1}}$$

where we keep B-smooth $g^{x_i} \bmod p \approx O(p)$ until we get more equations than primes p_i.

If we take log base g on both sides: $x_1 \equiv \sum e_{i,1} \log_g p_i \pmod{p-1}$.[3] Since we know the x_i and $e_{i,j}$, we can solve the system of linear equations for the discrete logs $\log_g p_i$ (since there are at least $t+1$ equations and t variables).

Now, take random y find an $h^y = p_1^{f_1} \dots p_t^{f_t}$ that is B-smooth. Taking logs as above, $y\alpha = \sum f_i \log_g p_i$ and we can solve for α.

Since this is basically the same process as random squares, it is no surprise it has similar time complexity $L_p(\frac{1}{2}, \sqrt{2})$. Practically, it's slightly harder than factoring.

---

[3]$\log_g p_i$ always exists. If g is not a generator since *some* generator h exists and we have $\frac{\log_h p_i \bmod p-1}{\log_h g \bmod p-1}$.

Chapter 4

Signatures

Lecture 19   Hash Functions (10/28)

To establish something that is NM-CCA2 secure, we need to somehow "sign" the ciphertext to distinguish "authenticated" ciphertexts. We can do this with MACs (e.g., AES-GCM or ChaCha20-Poly1305) but we will do something different.

Consider a hybrid encryption scheme: use public-key encryption to send a symmetric key that encrypts the message. This is CO 487 content.

Hash functions

Most common hash functions are the SHA family: SHA0 (broken 2005), SHA1 (broken 2017), SHA2 (actually used), SHA3 (not really used, made in anticipation of SHA2 breaking). Again, CO 487 content beyond the scope of this course.

Definition (*hash function*)

> Function $H : S \to T$ (typically, $S = \{0,1\}^*$ and $T = \{0,1\}^\lambda$)

Ideally, a hash function is a random oracle, i.e., $H \overset{\$}{\leftarrow} \{f : (f : S \to T)\}$. This is useful, e.g., for making hashed RSA signatures existentially unforgeable under chosen message attack.

There is no way to easily construct a random oracle because (1) we can't construct the set of all functions and (2) we run into measure theory issues with defining a probability distribution on that set. Instead we construct with desired properties:

1. Preimage resistant: Given $t \in T$, it is infeasible to find $s \in H^{-1}(t)$.

2. Second preimage resistant: Given $s \in S$, it is infeasible to find $s' \in S$ such that $s \neq s'$ and $H(s) = H(s')$.

3. Collision resistant: It is infeasible to find $s \neq s'$ such that $H(s) = H(s')$.

Example 19.1. Are all preimage resistant functions second preimage resistant?

Solution. Consider $f(x) = x^2 \bmod n$. To find a preimage, take $x = \sqrt{y}$ (hard). To find a second preimage, take $x' = -x \neq x$ so $(-x)^2 = x^2$ (easy). □

To be formal, use games. For example, with collision resistance: Suppose we have a family of hash functions $HashGen : \mathbb{1}^\lambda \mapsto H_\lambda$. Play the game:

1. Pick a hash function $H_\lambda \xleftarrow{\$} HashGen(\mathbb{1}^\lambda)$

2. $(s, s') \leftarrow \mathcal{A}(\mathbb{1}^\lambda, H_\lambda)$

with win condition $H_\lambda(s) = H_\lambda(s')$ and $s \neq s'$. We define $\{H_\lambda : \lambda \in \mathbb{N}\}$ to be collision-resistant if no probabilistic polynomial time adversary $\mathcal{A}$ can win this game with non-negligible probability in λ.

We can construct collision-resistant hash functions from claw-free permutations by Damgård.

Definition (*claw-free permutation*)

> Given a set X, the pair of permutations (f, g) is claw-free if it is infeasible to find $x_1, x_2 \in X$ such that $f(x_1) = g(x_2)$.

The wrong way: Given claw-free permutations $f : X \to X$ and $g : X \to X$, we define $H : \{0, 1\}^* \to X$ with $H(\varepsilon) = x_\varepsilon$. Inductively, $H(b_1 b_2 \cdots b_n) = h(H(b_1 \cdots b_{n-1}))$ where $h = f$ if $b_n = 0$ and g if $b_n = 1$. Claim this is collision-resistant because if there is a collision $H(m) = H(m')$ and $m \neq m'$, we have a claw at some point, which is a contradiction. Unfortunately, we could run into a loop back to x_ε.

Instead, pick $x_0 \in X$ and define $x_\varepsilon = g(f(x_0))$ and define $H(b_0 \cdots b_n) = h(h(H(b_0 \cdots b_{n-1})))$ as above. Then, we cannot arrive at x_ε because generating pairs of $f(f(\dots))$ and $h(h(\dots))$ cannot create $g(f(\dots))$.

Lecture 20   Signature Schemes (10/31)

Consider some RSA modulus $n = pq$, $p > 2$, $q > 2$, $p \neq q$ where $p \equiv 3 \pmod 4$ and $q \equiv 3 \pmod 4$ (Blum integers, notable for use in the Blum–Blum–Shub generator).

Let y_p and $-y_p$ be square roots of $y \bmod p$ (and for q). Notice that $\left(\frac{-1}{p}\right) = -1$ and $\left(\frac{-1}{q}\right) = -1$. Then, exactly one of $\{y_p, -y_p\}$ is a square mod p (and for q).

Finally, combining gives exactly one of the square roots of $y \bmod n$ is a square mod n. This means that $f(x) = x^2$ is a permutation on $((\mathbb{Z}/n\mathbb{Z})^\times)^2$.

Choose $a \in (\mathbb{Z}/n\mathbb{Z})^\times$ such that $\left(\frac{a}{n}\right) = -1$ Then, define $g(y) = a^2 y^2$ which is also a permutation.

Note: suppose $f(x) = g(y)$. Then, $x^2 = (ay)^2$ and $x \neq \pm ay$ because if $x = \pm ay$ then $\left(\frac{x}{n}\right) = \left(\frac{\pm 1}{n}\right)\left(\frac{a}{n}\right)\left(\frac{y}{n}\right)$ but this is $1 = (1)(-1)(1)$, contradiction. From this, we can factor n (by Fermat).

Overall: claw-free permutations $\to$ collision-resistant hash functions $\to$ {secure digital signatures, CCA2-secure encryption, etc.}

How do we generate secure digital signatures?

Suppose we have RSA $pk = (n, e)$ and $sk = (n, d)$. Then, define signing and verification as $Sign : (\mathbb{Z}/n\mathbb{Z})^\times \to (\mathbb{Z}/n\mathbb{Z})^\times m \mapsto \sigma := m^d \bmod n$ and $Verify : (m, \sigma) \mapsto \sigma^e \bmod n \stackrel{?}{=} m$.

Definition (*signature schemes*)

A signature (scheme) is a tuple $(KeyGen, Sign, Verify)$ where
- $KeyGen : \mathbb{1}^\lambda \mapsto (pk, sk)$
- $Sign : (sk, m) \mapsto \sigma$
- $Verify : (pk, m, \sigma) \mapsto \{0, 1\}$

and we have that if $(pk, sk) \xleftarrow{\$} KG(\mathbb{1}^\lambda)$ and $\sigma \xleftarrow{\$} S(sk, m)$, then $V(pk, m, \sigma) = 1$.

Under Textbook RSA, it is trivial to forge junk (but valid) signatures, i.e., given random signature σ, it signs some calculable message.

Example security definition game: EUF-CMA
Existential unforgeability (EUF): adversary produces a valid signature
Chosen-message attack (CMA): adversary can always use a signing oracle

1. $(pk, sk) \xleftarrow{\$} KeyGen(\mathbb{1}^\lambda)$
2. **for** $i = 1 \ldots q$ **do:**

   $m_i \xleftarrow{\$} \mathcal{A}(\mathbb{1}^\lambda, pk, (m_1, \sigma_1), \ldots, (m_{i-1}, \sigma_{i-1}))$

   $\sigma_i \xleftarrow{\$} Sign(sk, m_i)$

   end
3. $(m, \sigma) \xleftarrow{\$} \mathcal{A}(\mathbb{1}^\lambda, pk, (m_1, \sigma_1), \ldots, (m_q, \sigma_q))$

with win condition $Verify(pk, m, \sigma) = 1$ and for all i, $m \neq m_i$.

Definition (*EUF-CMA*)

A signature scheme is EUF-CMA if there does not exist a probabilistic polynomial time adversary $\mathcal{A}$ which wins the EUF-CMA game with non-negligible probability.

Hashed RSA $KeyGen : \mathbb{1}^\lambda \mapsto ((n, e), (n, d))$
$Sign : m \mapsto H(m)^d \bmod n$ for hash $H : \{0, 1\}^* \to (\mathbb{Z}/n\mathbb{Z})^\times$ (i.e., a claw-free permutation)
$Verify : (m, \sigma) \mapsto H(m) \overset{?}{=} \sigma^e \bmod n$

We can prove that if the RSA assumption holds[1] and the hash function H is a random oracle, then Hashed RSA is EUF-CMA.

Lecture 21   Hashed RSA (11/02)

Recall EUF-CMA and Hashed RSA. We want to prove

Theorem

Hashed RSA is EUF-CMA assuming:
- The RSA assumption holds
- The hash functions H are random oracles

Proof. For a contradiction, let $\mathcal{A}$ be an adversary that wins the EUF-CMA game, generating a forged signature (m_*, σ_*). Note that we must expose the hash function H to the adversary.

Consider when H has the property that for some $\sigma \in (\mathbb{Z}/n\mathbb{Z})^\times$, $H(m_*) = m\sigma^e$. Then, $\sigma_* = H(m_*)^d = (m\sigma^e)^d = m^d\sigma$ so $\sigma_*\sigma^{-1} = m^d$. We could return $H(m_*) = m\sigma^e$ but we

---

[1]Given n, e, m^e, it is infeasible to find m.

still have to respond to signing queries of $\mathcal{A}$ somehow.

To respond to a query for m_i, pick a random $\sigma_i \xleftarrow{\$} (\mathbb{Z}/n\mathbb{Z})^\times$ and set $H(m_i) = \sigma_i^e$ and respond with σ_i. Note that the challenger must maintain a table of $H(m_i)$ to respond to duplicates.

To make this work somehow, we define $H(m) = \begin{cases} m\sigma^e & \text{with probability } \frac{1}{q+1} \\ \sigma^e & \text{with probability } \frac{q}{q+1} \end{cases}$.

Then, notice that the adversary will make at most $q + 1$ relevant hash function requests (q for the signing queries, 1 for m_*). Now, the probability that we get what we want, i.e., calculate m^d, is $\left(\frac{q}{q+1}\right)^q \frac{1}{1+1} Adv(\mathcal{A}) > \frac{1}{(q+1)\exp(1)} Adv(\mathcal{A})$ which is non-negligible since q is polynomial and $Adv(\mathcal{A})$ is non-negligible.

That is, we can break RSA in probabilistic polynomial time with non-negligible probability, violating the RSA assumption. Therefore, $\mathcal{A}$ cannot exist. $\qquad\square$

Note: non-negligible means that there exists an n such that $\Pr[\mathcal{A} \text{ wins}] = f(\lambda) \in \Omega(\frac{1}{\lambda^n})$.

Further reading: EdDSA (Schnorr), "Short signatures without random oracles" (Boneh–Boyen)

Lecture 22   Zero-Knowledge Proofs (11/04)

Suppose that $x \in (\mathbb{Z}/n\mathbb{Z})^\times$ where $n = pq$.

Claim: there exists a y such that $x = y^2 \bmod n$.

If Alice knows that $x = y^2$ and sends y to Bob, that is a full-knowledge proof. A zero-knowledge proof would not send y.

Instead, Alice chooses a random $r \xleftarrow{\$} (\mathbb{Z}/n\mathbb{Z})^\times$ and computes $xr^2 = y^2r^2 = (yr)^2$. If she sends $\beta = xr^2$ and $\alpha = yr$, Bob can verify that $\alpha^2 = \beta$. However, Bob cannot trust that α is in fact yr and cannot prove that $\frac{\beta}{x} = r^2$ without sending r.

Protocol   For Alice to prove that she knows y such that $y^2 = x$,

1. Alice picks $r \xleftarrow{\$} (\mathbb{Z}/n\mathbb{Z})^\times$ and sends xr^2
2. Bob picks $b \xleftarrow{\$} \{0, 1\}$ and sends b
3. Alice sends $\rho = y^b r$ and sends ρ
4. Bob verifies that $\rho^2 = \beta x^{b-1}$

Then, if $b = 0$, Bob can catch a forged y and if $b = 1$, Bob is more certain that y exists. The chance that Alice is cheating and avoids being caught in λ iterations is $2^{-\lambda}$.

Suppose Alice does not know a square root $y = \sqrt{x}$. She could:

- Choose r randomly, send $\beta = xr^2$, and hope that $b = 0$ to send r
- Choose α randomly, send $\beta = \alpha^2$, and hope that $b = 1$ to send α

meaning that Alice can forge with success probability 50%, and indeed Bob could fool himself half the time by doing this himself. That is, a zero-knowledge proof does not introduce any new information that Bob could not have produced on his own.

Then, a security definition for a ZKP protocol requires

1. Correctness: With an honest prover and honest verifier, the proof succeeds with probability 100%.

2. Soundness: With a dishonest prover and honest verifier, then there is a non-negligible probability that they get caught.

3. Zero-knowledge: With an honest prover and dishonest verifier, then the verifier can simulate correct proofs with non-negligible probability. This means that the verifier cannot actually use any information for anything else (e.g., cannot factor a number even if the ZKP proves that the prover knows the factors).

where non-negligible means any useful number (e.g., 50%, 25%, 30%, etc.).

From a ZKP, we can construct a signature scheme. Generate a key (x, y) where $y^2 = x$. Signing is done by:

1. Alice picks random r and sends x and $\beta = xr^2$.
2. Bob picks random b and sends b
3. Alice calculates $\rho = ry^b$ and sends ρ

To verify, ensure that $\rho^2 = \beta x^{-1}$.

Alternatively, if we want to use DLOG (i.e., with g and g^x, prove that you know x):

1. Alice picks random r and sends g, g^x, and $\beta = g^r$
2. Bob picks random bit b and sends b
3. Alice calculates $\rho = r + bx$ and sends ρ

To verify, ensure that $g^\rho = \beta(g^x)^b$.

We call these Σ protocols because the back-and-forth looks like a Σ.

Definition (*Fiat–Shamir transformation*)

> To transform a ZKP protocol to a signing scheme, set $b = H(\beta, m)$ where H is a random oracle. Then, the signature of m is (β, ρ). To verify, assert ρ satisfies the ZKP protocol given β.

Notice that there is only one bit of entropy, so it is forgeable 50% of the time. If we try increasing entropy in the DLOG scheme by making $b \in \mathbb{Z}$, correctness and soundness still hold but zero-knowledge might not.

Lecture 23   ZKP Signatures (11/07)

Recall the Fiat–Shamir transform:

Given a Σ protocol, Peggy sends Victor the problem π and a commitment c (usually some sort of randomized value). Victor returns a challenge b. Peggy's response r depends on b.

From the perspective of a signature scheme, we generate a private key (statement to be proved) and public key π.

To sign, choose a commitment at random $c \xleftarrow{\$} C$. Set the challenge to be a deterministic but random function $b = H(m, c)$. Calculate a response for b. Then, let $\sigma = (c, r)$.

To verify, recalculate the challenge and verify with the response.

Notice that if $b \in \{0, 1\}$, then signature forgery is permitted half the time. To actually use this, the challenge space must be large.

Generically, the signer repeats the protocol λ times and initially commit to a commitment vector $\mathbf{c} = (c_1, \ldots, c_\lambda)$. They also make a challenge vector $\mathbf{b} = (b_1, \ldots, b_\lambda) = H(m, \mathbf{c})$. Then, to successfully fake the proof (and forge a signature), the signer would need to very luckily get a $\mathbf{b}$ that matches perfectly with a malicious $\mathbf{c}$. Finally, generate a response vector $\mathbf{r}$ and return $\sigma = (\mathbf{c}, \mathbf{r})$.

This cannot be proved to be ZK because the proof cannot be simulated. We assume that the heuristic (that ZKP's produce ZKP's) holds.

Example 23.1. Why do we need to generate the vectors at once?

WLOG, suppose we want to prove knowledge of x (in the DLOG problem).

The public key is $\pi = g^x$. The commitment is $c = g^y$ for random y. The challenge is a bit b. The response is $r = y + bx$.

If Peggy predicts $b = 0$, pick y randomly, set $c = g^y$, and Victor verifies $r = y$. Otherwise, if she predicts $b = 1$, pick r_0 randomly, set $c = \frac{g^{r_0}}{g^x}$, and Victor verifies $r = r_0$.

If Peggy continuously generates random y, she only has to try twice until getting desired $b = 0$. Then, she only needs to do 2λ work instead of 2^λ work.

What we're describing is the Schnorr signature.

Schnorr scheme   Given a group G and $g \in G$, generate keys $(pk, sk) = (g^x, x)$.

A signature of m is a proof of knowledge of x. First, generate a commitment $r \xleftarrow{\$} \mathbb{Z}$, $c = g^r$. Then, calculate a challenge $b \leftarrow H(m, c) = H(m, g^r)$. The response is $r + bx$, so return $(c, \sigma) = (g^r, r + bx)$.

To verify a signature (c, σ), check if $g^\sigma \stackrel{?}{=} c(g^x)^{H(m,g^r)}$. That is, compute $b' \leftarrow H(m, c) = H(m, g^r)$ and check if $g^{r+bx} = g^r(g^x)^b \stackrel{?}{=} c(pk)^{b'}$.

Alternatively, send (b, σ). Then, calculate $c' \leftarrow \frac{g^\sigma}{(g^x)^b}$ and check if $H(m, c') \stackrel{?}{=} b$. This is better since b is an integer, which is easier to serialize and send than a group element c.

Theorem (*Schnorr*)

> Assuming that H is a random oracle and that DLOG is hard in G, the Schnorr signature scheme is EUF-CMA.

Proof. Suppose we are an adversary $\mathcal{A}_1$ trying to solve DLOG. Let g, g^α be a challenge from $\mathcal{C}_1$. Suppose also that we are a challenger $\mathcal{C}_2$ with access to an adversary $\mathcal{A}_2$ that breaks Schnorr.

Give $\mathcal{A}_2$ the parameter g^α. Then, the adversary forges a signature $(b, r + b\alpha)$. We want to isolate α, so we need two signatures with the same public key, same commitment, but with different hashes b and b'. Using the *forking lemma*, we stop execution before the hashing and swap out the hash function H.

Then, we have $(b, r + b\alpha)$ and $(b', r + b'\alpha)$. We can now solve for α and return it to $\mathcal{C}_1$.

Since $\mathcal{A}_2$ runs in poly. time, we ($\mathcal{A}_1$) ran in poly. time, meaning that DLOG is easy. □

Lecture 24   CCA2-Secure Signature Schemes (11/09)

Recall: in IND-CCA2, $\mathcal{A}$ can use a decryption oracle, then produce two messages. $\mathcal{C}$ picks a random one of the two and encrypts it. Then, $\mathcal{A}$ gets access to the decryption oracle and wins if they can distinguish which message was encrypted.

In Textbook RSA, $E(m) = m^e \bmod n$, so an attacker can pick garbage k and ask for the decryption of $m^e k^e$. The core issue here is that E is a group homomorphism, i.e., $E(m_1 m_2) = E(m_1) E(m_2)$.

Remark. Any homomorphic cryptosystem is not CCA2-secure.

For example, Rabin encryption $E(m) = m^2 \bmod n$ is homomorphic and Elgamal $E(m) = (g^y, g^{xy} m)$ is also homomorphic in each entry.

Symmetric + asymmetric hybrid   Let $KeyGen$, $Enc : M \to C$, and $Dec : C \to M$ be a public key cryptosystem. Also let $\mathcal{E}nc$ and $\mathcal{D}ec$ be a symmetric key cryptosystem.

Suppose Alice wants to send to Bob. Bob generates $(pk, sk) \xleftarrow{\$} KeyGen$ and publishes pk.

Alice picks random $\sigma \xleftarrow{\$} M$, encrypts both $c = Enc(pk, \sigma)$ and $d = \mathcal{E}nc(\sigma, m)$, and sends (c, d). Notice that we can reinterpret σ as a key for the SKC by just treating it as an appropriately-sized bistring.

Bob can now decrypt (c, d) by first decrypting $\sigma = Dec(sk, c)$ and then $m = \mathcal{D}ec(\sigma, d)$.

Fujisaki–Okamoto   (1999) is a CCA2-secure one-time pad (OTP) hybrid. Let $KeyGen$, Enc, Dec be a PKC. Then, make a pseudo-OTP $\mathcal{E}nc(k, m) = m \oplus H_1(k)$ and add a MAC $H_2(k, m)$.

Generate $(pk, sk) \xleftarrow{\$} KeyGen(\mathbb{1}^\ell)$ and pick $\sigma \xleftarrow{\$} M$.

Then, $E(m) = (Enc(pk, \sigma), m \oplus H_1(\sigma), H_2(\sigma, m))$.

To invert, $D(c, d, e) = d \oplus H_1(Dec(sk, c)) = m$ and check $H_2(\sigma, m) = e$. If the MAC does not check out, either explicitly error or implicitly output random garbage $H_3(s, (c, d, e))$ with a secret seed s.

Then, the CCA2 oracle is sabotaged.

Lecture 25   Proving Fujisaki–Okamoto Security (11/11)

Recall the Fujisaki–Okamoto inputs: $KGen : \mathbb{1}^\ell \mapsto (pk, sk)$, $Enc : M \to C$, $Dec : C \to M$, $H_1 : M \to \{0, 1\}^n$, and $H_2 : \{0, 1\}^n \times M \to T$.

Then, $\mathcal{E}nc(pk, m) = (Enc(pk, r), m \oplus H_1(\sigma), H_2(m, \sigma))$ where $m \in \{0, 1\}^n$ and $\sigma \xleftarrow{\$} M$.

Theorem

> If the original PKC is OW-CPA and H_1, H_2 are reandom oracles, then this basic Fujisaki–Okamoto is IND-CPA.

Proof. Let $\mathcal{A}$ be an adversary that can win IND-CPA for FO. Recall IND-CPA: let $m_0, m_1 \leftarrow \mathcal{A}(\mathbb{1}^\ell, pk)$ and $b' \leftarrow \mathcal{A}(\mathbb{1}^\ell, pk, \mathcal{E}nc(pk, m_b))$. Then, $\mathcal{A}$ can find $b = b'$ with

non-negligible probability.

Notice that the second term $m \oplus H_1(\sigma)$ is garbage since σ is random so m is randomly scrambled. Therefore, it is information-theoretically indistinguishable from random garbage. So the only way to get any information about m is to find σ.

Therefore, $\Pr[\mathcal{A} \text{ wins}] \leq \Pr[\mathcal{A} \text{ finds } \sigma]$.

Suppose we are challenged to break the PKC in the OW-CPA game and are given (pk, σ).

Then, we can challenge $\mathcal{A}$ with $(Enc(pk, \sigma), \tau, \mu)$ with random garbage τ, μ. Then, at some point $\mathcal{A}$ must call $H_1(\sigma)$. We intercept all the calls to H_1 (since we control H_1) and respond with σ with non-negligible probability. Therefore, if FO is not IND-CPA, then the PKC is not OW-CPA. $\qquad\square$

This reduction is not *tight* because we randomly pick potential σ candidates. If the original PKC is deterministic, then we can re-encrypt all potential σ to find the right one.

Theorem

> If Enc is deterministic, the PKS is OW-CPA, and H_1, H_2 are random oracles, then FO is IND-CCA2.

Proof. First, notice that the IND-CCA2 game without the decryption oracle is the IND-CPA game.

However, in FO, we claim the decryption oracle is "useless" because there is no information-theoretic use of it. Therefore, since FO is IND-CPA, it is also IND-CCA2.

To prove the claim, consider $\mathcal{E}nc(pk, m) = (Enc(pk, \sigma), m \oplus H_1(\sigma), H_2(m, \sigma))$.

Then, $\mathcal{D}ec(sk, (c_1, c_2, c_3)) = \begin{cases} \underbrace{H_1(\overbrace{Dec(sk, c_1)}^{\sigma'}) \oplus c_2}_{m'} & \text{otherwise} \\ \bot & c_3 \neq H_2(m', \sigma') \end{cases}$

Since encryption is deterministic, the only way to construct a valid ciphertext that gets a return value is to know both m and σ to calculate $Enc(pk, m)$ and $H_2(m, \sigma)$.

We can simulate this for the adversary by intercepting calls to H_2 and checking if σ matches the encryption of m (i.e. c_1). Therefore, there is no difference between IND-CPA and IND-CCA2. $\qquad\square$

Full Fujisaki–Okamoto   Instead of using $Enc(pk, \sigma)$, randomize to $Enc(pk, \sigma; r)$. For example, in Elgamal, $Enc(g^x, \sigma; r) = (g^r, g^{xr}\sigma)$.

Then, $\mathcal{E}nc(pk, m) = (Enc(pk, \sigma; H_2(m, \sigma)), m \oplus H_1(\sigma))$ for $\sigma \xleftarrow{\$} M$. That is, we use the tag as the randomness.

Finally, $\mathcal{D}ec(sk, (c_1, c_2)) = \begin{cases} \underbrace{H_1(\overbrace{Dec(sk, c_1)}^{\sigma'}) \oplus c_2}_{m'} & \text{otherwise} \\ \bot & c_1 \neq Enc(pk, \sigma'; H_2(m', s')) \end{cases}$

Theorem

> If the PKC is OW-CPA and H_1, H_2 are random oracles, then full FO is IND-CCA2.

What if we don't have random oracles? Cramer–Shoup (1998) gets IND-CCA2 using DDH

and a collision-resistant hash function. It is also stupid complicated.

Given a group $|G| = q$ with two generators g_1, g_2 where $\langle g_1 \rangle = \langle g_2 \rangle = G$.

The $sk = (x_1, x_2, y_1, y_2, z) \in (\mathbb{Z}/q\mathbb{Z})^5$ and $pk = (c, d, h) = (g_1^{x_1} g_2^{x_2}, g_1^{y_1} g_2^{y_2}, g_1^z)$.

Encryption is $Enc(pk, m) = (g_1^r, g_2^r, h^r m, c^r d^{r H(g_1^m, g_2^m, h^r m)})$ for $m \in G$ and $r \xleftarrow{\$} \mathbb{Z}/q\mathbb{Z}$.

Then, the last part $c^r d^{r\alpha}$ acts as a checksum. To generate a valid ciphertext and use the CCA2 oracle, an adversary must generate this, which breaks DDH.

Chapter 5

Elliptic Curve Cryptography

Lecture 26 Elliptic Curves (11/14)

Recall the conic sections: $y^2 = 1 - x^2$ (circles), $y^2 = x^2 - 1$ (hyperbola), etc. If we replace the quadratic in x with a cubic, we get an elliptic curve. For our purposes,

Definition (*curve*)

> The set of points satisfying $f(x, y) = 0$ where $f \in K[x, y]$ for a field K.

where K is a (usually finite) field, e.g., $\mathbb{Z}/p\mathbb{Z}$ or something funny like $\mathbb{Z}/3\mathbb{Z}[i] = \mathbb{F}_9$.

Note that we can rewrite any cubic $ax^3 + bx^2 + cx + d$ by first dividing through by a to get $x^3 + b'x^2 + c'x + d'$. Then, send $x \mapsto x - \frac{b'}{3}$ to get $x^3 + c''x + d''$. This only works if $3 \neq 0$ so we can divide by 3, i.e., the characteristic of K is not 3.

To simplify the quadratic in y, we can complete the square as long as $2 \neq 0$, i.e., the characteristic of K is not 3.

Definition (*elliptic curve*)

> A solution set to an equation of the form $y^2 = x^3 + ax + b$ where $a, b \in K$ and $\mathrm{char}(K) \neq 2, 3$.

Consider an ellipse centered at the origin with semimajor axes a and b. Then, the arc length is $\int_0^{\pi/2} \sqrt{a^2 \cos^2 t + b^2 \sin^2 t}\,\mathrm{d}t$. Make the substitution $u = \sin t$, $\mathrm{d}u = \cos t\,\mathrm{d}t$ to get $\int \sqrt{a^2 - \frac{(a^2 - b^2)u^2}{1 - u^2}}\,\mathrm{d}u$. Then, $k^2 = 1 - \frac{b^2}{a^2}$ gives $\int a\sqrt{\frac{1 - k^2 u^2}{1 - u^2}}\,\mathrm{d}u$ and finally $x = 1 - k^2 u^2$ for $\frac{1}{2}\int_{1-k^2}^1 \frac{x\,\mathrm{d}x}{\sqrt{x(x-1)(x-(1-k^2))}}$. This is our *elliptic integral.*

Generally, elliptic integrals of the first kind $\int \frac{\mathrm{d}x}{\sqrt{x^3 + \cdots}}$ and of the second kind $\int \frac{x\,\mathrm{d}x}{\sqrt{x^3 + \cdots}}$.

Just as $\int \frac{\mathrm{d}x}{\sqrt{x^2 + \cdots}}$ gives $\sin^{-1}(x)$, the inverse of a periodic function, complex analysis shows that elliptic integrals of the first kind gives the inverse of the doubly periodic Weierstrass function $\wp^{-1}(x)$.

By analogy to circles which can be defined by $f^2 + f'^2 = 1$, elliptic integrals of the first kind satisfy $\wp'^2 = 4\wp^3 + c_1\wp + c_2$.

Elliptic curves have a somewhat natural group law.

Lemma

> Every line intersects an elliptic curve in exactly three places (up to multiplicity).

Proof. Let $y^2 = x^3 + ax + b$ be an elliptic curve.

Consider a line L through $P = (x_P, y_P)$ and $Q = (x_Q, y_Q)$. Then, the slope of L is $\frac{y_Q - y_P}{x_Q - x_P} = m$. Substitute $y = mx + c$ into the elliptic curve to get a cubic in x. The cubic has three roots.

Then, it is $(x - x_P)(x - x_Q)(x - x_R)$. Taking the coefficient on x^2, we have $x_R = m^2 - x_P - x_Q$ and $y_R = m((m^2 - x_P - x_Q) - x_P) + y_P$. $\qquad\square$

Define the group law as $P + Q = (x_R, -y_R)$.

What is $P + P$? Take the tangent line to P, i.e., "$\lim_{Q \to P}(P + Q)$".

What if there is no tangent line (self-intersection or cusps)? To ensure this cannot happen, assert that the discriminant $4a^3 + 27b^2 \neq 0$.

What is the identity? Not really a point, denote ∞ or $[0 : 1 : 0]$ in Sage.